

```
In [1]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

from sklearn.model_selection import cross_val_score, GridSearchCV, train_test_split
from sklearn.feature_selection import SelectFromModel

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier

import matplotlib.pyplot as plt
```

```
In [2]: #Import soccer data(database.sqlite) using pandas.
import pandas as pd
import sqlite3
db = 'database.sqlite'
conn = sqlite3.connect(db)
cur = conn.cursor()
```

```
In [3]: #Load the values from the attributes X= ['strength', 'stamina'] and Y= ['jumping'] from

df = pd.read_sql("SELECT strength, stamina, jumping FROM Player_Attributes", conn)
df.fillna(11, inplace=True)

df.head()
```

```
Out[3]:
```

	strength	stamina	jumping
0	76.0	54.0	58.0
1	76.0	54.0	58.0
2	76.0	54.0	58.0
3	76.0	54.0	58.0
4	76.0	54.0	58.0

```
In [4]: x = df[['strength', 'stamina']].values
y = df[['jumping']].values
X_train, X_test, y_train, y_test = train_test_split(x, y)
```

```
In [5]: #Using the above X and Y, apply GridSearch on DecisionTreeClassifier and fetch the best
from sklearn import metrics

model = DecisionTreeClassifier()
dt_mod = model.fit(X_train, y_train)
dt_preds = dt_mod.predict(X_test)
accuracy = metrics.accuracy_score(y_test, dt_preds)
accuracy
```

```
Out[5]: 0.16460484835308187
```

```
In [12]: #apply gridsearch

model.get_params()
```

```
Out[12]: {'ccp_alpha': 0.0,
          'class_weight': None,
          'criterion': 'gini',
          'max_depth': None,
          'max_features': None,
          'max_leaf_nodes': None,
          'min_impurity_decrease': 0.0,
          'min_impurity_split': None,
          'min_samples_leaf': 1,
          'min_samples_split': 2,
          'min_weight_fraction_leaf': 0.0,
          'random_state': None,
          'splitter': 'best'}
```

```
In [30]: #params_grid = {'criterion':['gini','entropy'],'splitter':['best','random'],'max_depth'

params_grid = {'criterion':['gini','entropy'], 'max_depth':[4,5,6,7,8,9,10,11,12,15,20,
```

```
In [31]: grid_search = GridSearchCV(DecisionTreeClassifier(), params_grid)
```

```
In [32]: grid_search.fit(X_train, y_train)
```

C:\Users\benso\Anaconda3\lib\site-packages\sklearn\model_selection_split.py:668: UserWarning: The least populated class in y has only 1 members, which is less than n_splits=5. % (min_groups, self.n_splits)), UserWarning)

```
Out[32]: GridSearchCV(estimator=DecisionTreeClassifier(),
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_depth': [4, 5, 6, 7, 8, 9, 10, 11, 12, 15, 20, 30,
                                                40, 50],
                                   'random_state': [0, 1, 2, 4, 6, 8, 10, 12, 14, 16, 20,
                                                40, 42],
                                   'splitter': ['best', 'random']})
```

```
In [34]: #Using this best params (from 3.1) and the selected features (from 3.3), rebuild the De

print("best score: " + str(grid_search.best_score_))
print("test score: " + str(grid_search.score(X_test, y_test)))
decision_tree_best_params = grid_search.best_params_
print("best parameters: " + str(decision_tree_best_params))
```

```
best score: 0.1596428582289639
test score: 0.16462658984672246
best parameters: {'criterion': 'entropy', 'max_depth': 20, 'random_state': 4, 'splitte
r': 'best'}
```

```
In [35]: #re-do with best parameters

model = DecisionTreeClassifier(criterion='gini',max_depth=20,random_state=4,splitter='b
dt_mod = model.fit(X_train, y_train)
dt_preds = dt_mod.predict(X_test)
accuracy = metrics.accuracy_score(y_test, dt_preds)
accuracy
```

```
#same as before....
```

Out[35]: 0.161169692357865

```
In [6]: #repeat for SVM
from sklearn import svm
from sklearn.svm import LinearSVC

#from sklearn.preprocessing import MinMaxScaler
#scaling = MinMaxScaler(feature_range=(-1,1)).fit(X_train)
#X_train = scaling.transform(X_train)
#X_test = scaling.transform(X_test)

from sklearn import preprocessing
X_train = preprocessing.scale(X_train)
X_test = preprocessing.scale(X_test)

model = LinearSVC(random_state=0, tol=1e-5)
svm_mod = model.fit(X_train, y_train.ravel())

#model = svm.SVC()
#dt_mod = model.fit(X_train, y_train.ravel())
svm_preds = svm_mod.predict(X_test)
accuracy = metrics.accuracy_score(y_test, svm_preds)
accuracy
```

```
In [8]: #repeat for LogisticRegression

from sklearn import metrics
from sklearn import preprocessing
X_train = preprocessing.scale(X_train)
X_test = preprocessing.scale(X_test)

model = LogisticRegression()
dt_mod = model.fit(X_train, y_train.ravel())
dt_preds = dt_mod.predict(X_test)
accuracy = metrics.accuracy_score(y_test, dt_preds)
accuracy
```

C:\Users\benso\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:765: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Out[8]: 0.053223176432220895

```
In [9]: model.get_params()
```

Out[9]: {'C': 1.0,
 'class_weight': None,
 'dual': False,

```
'fit_intercept': True,
'intercept_scaling': 1,
'l1_ratio': None,
'max_iter': 100,
'multi_class': 'auto',
'n_jobs': None,
'penalty': 'l2',
'random_state': None,
'solver': 'lbfgs',
'tol': 0.0001,
'verbose': 0,
'warm_start': False}
```

```
In [17]: params_grid = {'C':[1.0,2.0,3.0], 'max_iter':[100,200,300]}
```

```
In [18]: grid_search = GridSearchCV(LogisticRegression(), params_grid)
```

```
In [19]: grid_search.fit(X_train, y_train.ravel())
```

```
C:\Users\benso\Anaconda3\lib\site-packages\sklearn\model_selection\_split.py:668: UserWarning: The least populated class in y has only 1 members, which is less than n_splits=5.
  % (min_groups, self.n_splits)), UserWarning)
C:\Users\benso\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  return f(*args, **kwargs)
C:\Users\benso\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:765: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
C:\Users\benso\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  return f(*args, **kwargs)
```

KeyboardInterrupt Traceback (most recent call last)

<ipython-input-19-231b374c92ba> in <module>

```
----> 1 grid_search.fit(X_train, y_train)
```

```
~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inner_f(*args, **kwargs)
```

```
61         extra_args = len(args) - len(all_args)
62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
64
65         # extra_args > 0
```

```
~\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py in fit(self, X, y, groups, **fit_params)
```

```
839         return results
840
--> 841         self._run_search(evaluate_candidates)
842
843         # multimetric is determined here because in the case of a callable
```

```
~\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py in _run_search(self, ev
```

```

evaluate_candidates)
1286     def _run_search(self, evaluate_candidates):
1287         """Search all candidates in param_grid"""
-> 1288         evaluate_candidates(ParameterGrid(self.param_grid))
1289
1290

~\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py in evaluate_candidates
(candidate_params, cv, more_results)
807         (split_idx, (train, test)) in product(
808             enumerate(candidate_params),
--> 809             enumerate(cv.split(X, y, groups))))
810
811         if len(out) < 1:

~\Anaconda3\lib\site-packages\joblib\parallel.py in __call__(self, iterable)
1042         self._iterating = self._original_iterator is not None
1043
-> 1044         while self.dispatch_one_batch(iterator):
1045             pass
1046

~\Anaconda3\lib\site-packages\joblib\parallel.py in dispatch_one_batch(self, iterator)
857         return False
858     else:
--> 859         self._dispatch(tasks)
860         return True
861

~\Anaconda3\lib\site-packages\joblib\parallel.py in _dispatch(self, batch)
775         with self._lock:
776             job_idx = len(self._jobs)
-> 777             job = self._backend.apply_async(batch, callback=cb)
778             # A job can complete so quickly than its callback is
779             # called before we get here, causing self._jobs to

~\Anaconda3\lib\site-packages\joblib\_parallel_backends.py in apply_async(self, func, callback)
206     def apply_async(self, func, callback=None):
207         """Schedule a func to be run"""
--> 208         result = ImmediateResult(func)
209         if callback:
210             callback(result)

~\Anaconda3\lib\site-packages\joblib\_parallel_backends.py in __init__(self, batch)
570         # Don't delay the application, to avoid keeping the input
571         # arguments in memory
--> 572         self.results = batch()
573
574     def get(self):

~\Anaconda3\lib\site-packages\joblib\parallel.py in __call__(self)
261         with parallel_backend(self._backend, n_jobs=self._n_jobs):
262             return [func(*args, **kwargs)
--> 263                     for func, args, kwargs in self.items]
264
265     def __reduce__(self):

~\Anaconda3\lib\site-packages\joblib\parallel.py in <listcomp>(.0)
261         with parallel_backend(self._backend, n_jobs=self._n_jobs):
262             return [func(*args, **kwargs)
--> 263                     for func, args, kwargs in self.items]
264
265     def __reduce__(self):

```

```

~\Anaconda3\lib\site-packages\sklearn\utils\fixes.py in __call__(self, *args, **kwargs)
    220     def __call__(self, *args, **kwargs):
    221         with config_context(**self.config):
--> 222             return self.function(*args, **kwargs)

~\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py in _fit_and_score(estimator, X, y, scorer, train, test, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, return_estimator, split_progress, candidate_progress, error_score)
    591         estimator.fit(X_train, **fit_params)
    592     else:
--> 593         estimator.fit(X_train, y_train, **fit_params)
    594
    595     except Exception as e:

~\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py in fit(self, X, y, sample_weight)
    1414         penalty=penalty, max_squared_sum=max_squared_sum,
    1415         sample_weight=sample_weight)
-> 1416         for class_, warm_start_coef_ in zip(classes_, warm_start_coef))
    1417
    1418         fold_coefs_, _, n_iter_ = zip(*fold_coefs_)

~\Anaconda3\lib\site-packages\joblib\parallel.py in __call__(self, iterable)
    1039         # remaining jobs.
    1040         self._iterating = False
-> 1041         if self.dispatch_one_batch(iterator):
    1042             self._iterating = self._original_iterator is not None
    1043

~\Anaconda3\lib\site-packages\joblib\parallel.py in dispatch_one_batch(self, iterator)
    857         return False
    858     else:
--> 859         self._dispatch(tasks)
    860         return True
    861

~\Anaconda3\lib\site-packages\joblib\parallel.py in _dispatch(self, batch)
    775         with self._lock:
    776             job_idx = len(self._jobs)
--> 777             job = self._backend.apply_async(batch, callback=cb)
    778             # A job can complete so quickly than its callback is
    779             # called before we get here, causing self._jobs to

~\Anaconda3\lib\site-packages\joblib\_parallel_backends.py in apply_async(self, func, callback)
    206     def apply_async(self, func, callback=None):
    207         """Schedule a func to be run"""
--> 208         result = ImmediateResult(func)
    209         if callback:
    210             callback(result)

~\Anaconda3\lib\site-packages\joblib\_parallel_backends.py in __init__(self, batch)
    570         # Don't delay the application, to avoid keeping the input
    571         # arguments in memory
--> 572         self.results = batch()
    573
    574     def get(self):

~\Anaconda3\lib\site-packages\joblib\parallel.py in __call__(self)
    261         with parallel_backend(self._backend, n_jobs=self._n_jobs):
    262             return [func(*args, **kwargs)
--> 263                     for func, args, kwargs in self.items]
    264
    265     def _reduce__(self):

```

```

~\Anaconda3\lib\site-packages\joblib\parallel.py in <listcomp>(.0)
    261         with parallel_backend(self._backend, n_jobs=self._n_jobs):
    262             return [func(*args, **kwargs)
--> 263                     for func, args, kwargs in self.items]
    264
    265     def __reduce__(self):

~\Anaconda3\lib\site-packages\sklearn\utils\fixes.py in __call__(self, *args, **kwargs)
    220     def __call__(self, *args, **kwargs):
    221         with config_context(**self.config):
--> 222             return self.function(*args, **kwargs)

~\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py in _logistic_regression_
path(X, y, pos_class, Cs, fit_intercept, max_iter, tol, verbose, solver, coef, class_wi
ght, dual, penalty, intercept_scaling, multi_class, random_state, check_input, max_squar
ed_sum, sample_weight, l1_ratio)
    759         func, w0, method="L-BFGS-B", jac=True,
    760         args=(X, target, 1. / C, sample_weight),
--> 761         options={"iprint": iprint, "gtol": tol, "maxiter": max_iter}
    762     )
    763     n_iter_i = _check_optimize_result(

~\Anaconda3\lib\site-packages\scipy\optimize\_minimize.py in minimize(fun, x0, args, met
hod, jac, hess, hessp, bounds, constraints, tol, callback, options)
    618     elif meth == 'l-bfgs-b':
    619         return _minimize_lbfgsb(fun, x0, args, jac, bounds,
--> 620                                callback=callback, **options)
    621     elif meth == 'tnc':
    622         return _minimize_tnc(fun, x0, args, jac, bounds, callback=callback,

~\Anaconda3\lib\site-packages\scipy\optimize\lbfgsb.py in _minimize_lbfgsb(fun, x0, arg
s, jac, bounds, disp, maxcor, ftol, gtol, eps, maxfun, maxiter, iprint, callback, maxls,
finite_diff_rel_step, **unknown_options)
    358         # until the completion of the current minimization iteration.
    359         # Overwrite f and g:
--> 360         f, g = func_and_grad(x)
    361         elif task_str.startswith(b'NEW_X'):
    362             # new iteration

~\Anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in fun_and_g
rad(self, x)
    258         if not np.array_equal(x, self.x):
    259             self._update_x_impl(x)
--> 260         self._update_fun()
    261         self._update_grad()
    262         return self.f, self.g

~\Anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in _update_fun
(self)
    224     def _update_fun(self):
    225         if not self.f_updated:
--> 226             self._update_fun_impl()
    227             self.f_updated = True
    228

~\Anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in update_fun
()
    131
    132     def update_fun():
--> 133         self.f = fun_wrapped(self.x)
    134
    135         self._update_fun_impl = update_fun

~\Anaconda3\lib\site-packages\scipy\optimize\_differentiable_functions.py in fun_wrapped

```

```

(x)
128         def fun_wrapped(x):
129             self.nfev += 1
--> 130             return fun(x, *args)
131
132         def update_fun():

~\Anaconda3\lib\site-packages\scipy\optimize\optimize.py in __call__(self, x, *args)
72     def __call__(self, x, *args):
73         """ returns the the function value """
---> 74         self._compute_if_needed(x, *args)
75         return self._value
76

~\Anaconda3\lib\site-packages\scipy\optimize\optimize.py in _compute_if_needed(self, x,
*args)
66         if not np.all(x == self.x) or self._value is None or self.jac is None:
67             self.x = np.asarray(x).copy()
---> 68             fg = self.fun(x, *args)
69             self.jac = fg[1]
70             self._value = fg[0]

~\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py in func(x, *args)
734         target = Y_multi
735         if solver == 'lbfgs':
--> 736             def func(x, *args): return _multinomial_loss_grad(x, *args)[0:2]
737         elif solver == 'newton-cg':
738             def func(x, *args): return _multinomial_loss(x, *args)[0]

~\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py in _multinomial_loss_grad(w, X, Y, alpha, sample_weight)
346         grad = np.zeros((n_classes, n_features + bool(fit_intercept)),
347                          dtype=X.dtype)
--> 348         loss, p, w = _multinomial_loss(w, X, Y, alpha, sample_weight)
349         sample_weight = sample_weight[:, np.newaxis]
350         diff = sample_weight * (p - Y)

~\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py in _multinomial_loss(w,
X, Y, alpha, sample_weight)
297         p += intercept
298         p -= logsumexp(p, axis=1)[:, np.newaxis]
--> 299         loss = -(sample_weight * Y * p).sum()
300         loss += 0.5 * alpha * squared_norm(w)
301         p = np.exp(p, p)

~\Anaconda3\lib\site-packages\numpy\core\_methods.py in _sum(a, axis, dtype, out, keepdims, initial, where)
45 def _sum(a, axis=None, dtype=None, out=None, keepdims=False,
46          initial=_NoValue, where=True):
---> 47     return umr_sum(a, axis, dtype, out, keepdims, initial, where)
48
49 def _prod(a, axis=None, dtype=None, out=None, keepdims=False,

```

KeyboardInterrupt:

In [7]:

```

#repeat for KNeighborsClassifier

from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics

model = KNeighborsClassifier()

k_neigh = model.fit(X_train, y_train)

```



```
k_preds = k_neigh.predict(X_test)
accuracy = metrics.accuracy_score(y_test, k_preds)
accuracy
```

C:\Users\benso\Anaconda3\lib\site-packages\sklearn\neighbors_classification.py:179: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return self._fit(X, y)
```

Out[7]: 0.11312099141211002

In [8]: `model.get_params()`

```
Out[8]: {'algorithm': 'auto',
        'leaf_size': 30,
        'metric': 'minkowski',
        'metric_params': None,
        'n_jobs': None,
        'n_neighbors': 5,
        'p': 2,
        'weights': 'uniform'}
```

In [9]: `params_grid = {'algorithm': ['auto', 'ball_tree', 'kd_tree'], 'n_neighbors': [2, 5, 10, 20]}`

In [10]: `grid_search = GridSearchCV(KNeighborsClassifier(), params_grid)`

In [11]: `grid_search.fit(X_train, y_train.ravel())`

C:\Users\benso\Anaconda3\lib\site-packages\sklearn\model_selection_split.py:668: UserWarning: The least populated class in y has only 1 members, which is less than n_splits=5. % (min_groups, self.n_splits)), UserWarning)

```
Out[11]: GridSearchCV(estimator=KNeighborsClassifier(),
                      param_grid={'algorithm': ['auto', 'ball_tree', 'kd_tree'],
                                   'n_neighbors': [2, 5, 10, 20]})
```

In [12]: `print("best score: " + str(grid_search.best_score_))
print("test score: " + str(grid_search.score(X_test, y_test)))
K_neighbors_best_params = grid_search.best_params_
print("best parameters: " + str(K_neighbors_best_params))`

```
best score: 0.12534873906768934
```

```
test score: 0.12947059462985108
```

```
best parameters: {'algorithm': 'auto', 'n_neighbors': 20}
```

In []: *#two of the classifiers took more than 3 hours to run, so the data is inconclusive. Bas
'auto' algorithm, with 20 neighbors.*